Verdikt. · AI reads, deterministic code writes, humans decide

A decision pipeline for regulated document workflows. It reads the case, decides against the program's own rules, and sends a reply that **no language model touched**.

TRACK · BEST APPS & AGENTS

NEBIUS × NVIDIA GLOBAL AI HACKATHON 2026

HOW ONE CASE FLOWS · EVERY STEP BUILT AND RUNNING

HUMAN
Packet arrives
Email, JSON or MCP; fax and scans as attachments

›

IMAP · API
One case
Normalized, tenant-scoped, deduped

›

AI · VLM
Ingest reads
MiniCPM-V into the program's field schema

›

AI · LLM
Triage routes
Nemotron Nano, rule-based fallback

›

AI · LLM
Decide cites
Qwen3 retrieval, Nemotron Super verdict

›

NO AI
Deliver merges
Verdict into a frozen, approved template

›

HUMAN
Reviewer decides
Approve or reject; audit entry per stage

WHAT WE BUILT

AI-IN · REVIEWABLE

Models read and decide

- **Ingest** reads attachments into the program's own field schema
- **Triage** classifies the case and picks a route; refills skip the verdict
- **Retrieve** finds the rules that apply, by embedding
- **Verdict** is approve, deny or missing info, citing the rule that fired
- **Every model call** has a fallback and an audit entry

DETERMINISTIC-OUT

Code writes the reply

- **Frozen templates:** human-approved, versioned, hashed
- **Merged, not generated:** no model on the render path
- **Reply package:** the letter plus the enclosures that outcome needs
- **Boundary test** fails the build if render imports the AI layer
- **Reproducible:** template version and render fingerprint on record

PEOPLE AND CONTROL

Humans on exceptions

- **Pipeline view:** each stage's output, highlighted on the source PDF
- **Approve or reject** a held reply before it leaves
- **Rules library:** human-authored and versioned
- **Programs as data:** a new program needs no build
- **Tenant scoping,** roles, signed idempotent webhooks

GUARDRAILS The verdict **cites the rule that fired** · the template is **frozen before it can render** · **no model at render time** · every stage is audited with its model and time · one Nebius endpoint, **model per stage configurable**

WHAT IT PROVES

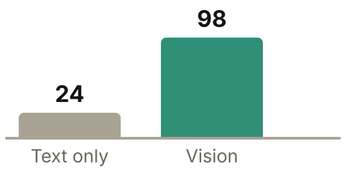
921
tests in the suite
916 passing, 0 failing, 5 skipped

0
models on render
Letter merged from a frozen template

5
stages, each observable
Ingest through Remember

2
programs, same code
One household passes one cap, fails the other

FIELDS RIGHT, OF 104



One dense form, one benchmark; not a guarantee

Pilot figures (re-open rate, time to verdict, cost per case) are still to be measured

A reusable pattern: any regulated letter that is also the audit artifact.

Test counts from the latest full run · demo data is fictional · no business-value figures are claimed